

```
In [16]: !jupyter nbconvert --to pdf AML_Neural_Network_Assignment_SriAnu(811362334).ipynb
```

```
/bin/bash: -c: line 1: syntax error near unexpected token `('
```

```
/bin/bash: -c: line 1: `jupyter nbconvert --to pdf AML_Neural_Network_Assignment_SriAnu(811362334).ipynb'
```

Neural Network Performance Analysis

In this notebook, I explore various neural network architectures to predict sentiment in the IMDB dataset. The models were trained with different combinations of hidden layers, hidden units, and loss functions. A detailed analysis is given below. [link text](#)

One-Hidden Layer Model

- **Overview:** A single hidden layer, comprised of 32 units with a `tanh` activation function, was trained on the model.
- **Performance:** Training accuracy improved slowly throughout the epochs, but fluctuated across several repeated epochs. Validation accuracy fluctuated around 50%, indicating possible underfitting.
- **Conclusion:** Using just one hidden layer was not very effective, probably due to the simplicity of the model being insufficient for this problem.

```
In [6]: # Import necessary modules
from tensorflow.keras import models, layers, regularizers
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
import matplotlib.pyplot as plt

# Load the IMDB dataset
(train_samples, train_labels), (test_samples, test_labels) = imdb.load_data(num_words=10000)

# Preprocess the data (pad sequences to make them the same length)
x_train_mod = pad_sequences(train_samples[:20000], maxlen=256)
y_train_mod = train_labels[:20000]
x_valid = pad_sequences(train_samples[20000:], maxlen=256)
y_valid = train_labels[20000:]

# Define a new model with one hidden layer
simple_nn_model = models.Sequential([
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(1, activation="sigmoid")
])

# Compile the model
simple_nn_model.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])


# Train the model
history_simple_nn = simple_nn_model.fit(x_train_mod, y_train_mod, epochs=20, batch_size=512, validation_data=(x_


# Function to plot training history with new variable names
def plot_new_history(history, graph_title):
    hist_dict = history.history
    loss_data = hist_dict["loss"]
    val_loss_data = hist_dict["val_loss"]
    accuracy_train = hist_dict.get("accuracy", None)
    accuracy_val = hist_dict.get("val_accuracy", None)
    epochs_range = range(1, len(loss_data) + 1)


    # Plot training and validation loss
    plt.plot(epochs_range, loss_data, "ro", label="Train Loss")
    plt.plot(epochs_range, val_loss_data, "r", label="Validation Loss")
    plt.title(f"Loss Analysis: {graph_title}")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.show()


    # If accuracy is available, plot training and validation accuracy
    if accuracy_train and accuracy_val:
        plt.clf()
        plt.plot(epochs_range, accuracy_train, "ro", label="Train Accuracy")
        plt.plot(epochs_range, accuracy_val, "r", label="Validation Accuracy")
        plt.title(f"Accuracy Analysis: {graph_title}")
        plt.xlabel("Epochs")
        plt.ylabel("Accuracy")
        plt.legend()
        plt.show()


# Plot the results for one hidden layer with new names
plot_new_history(history_simple_nn, "Modified One Hidden Layer Model")
```


Epoch 1/20
40/40  2s 13ms/step - accuracy: 0.5073 - loss: 0.8992 - val_accuracy: 0.5002 - val_loss: 0.8771


Epoch 2/20
40/40  0s 7ms/step - accuracy: 0.5053 - loss: 0.8663 - val_accuracy: 0.5018 - val_loss: 0.8439


Epoch 3/20
40/40  0s 6ms/step - accuracy: 0.5029 - loss: 0.8356 - val_accuracy: 0.4988 - val_loss: 0.8209


Epoch 4/20
40/40  0s 7ms/step - accuracy: 0.5019 - loss: 0.8117 - val_accuracy: 0.5006 - val_loss: 0.7964


Epoch 5/20
40/40  1s 6ms/step - accuracy: 0.5106 - loss: 0.7882 - val_accuracy: 0.4954 - val_loss: 0.7827


Epoch 6/20
40/40  0s 6ms/step - accuracy: 0.5075 - loss: 0.7736 - val_accuracy: 0.4984 - val_loss: 0.7685


Epoch 7/20
40/40  0s 6ms/step - accuracy: 0.5164 - loss: 0.7604 - val_accuracy: 0.4926 - val_loss: 0.7584


Epoch 8/20
40/40  0s 8ms/step - accuracy: 0.5143 - loss: 0.7509 - val_accuracy: 0.4950 - val_loss: 0.7513


Epoch 9/20
40/40  0s 6ms/step - accuracy: 0.5046 - loss: 0.7461 - val_accuracy: 0.4944 - val_loss: 0.7440


Epoch 10/20
40/40  0s 7ms/step - accuracy: 0.5130 - loss: 0.7390 - val_accuracy: 0.4964 - val_loss: 0.7383


Epoch 11/20
40/40  1s 6ms/step - accuracy: 0.5154 - loss: 0.7352 - val_accuracy: 0.5084 - val_loss: 0.7352


Epoch 12/20
40/40  0s 6ms/step - accuracy: 0.5199 - loss: 0.7316 - val_accuracy: 0.5030 - val_loss: 0.7335


Epoch 13/20
40/40  0s 7ms/step - accuracy: 0.5209 - loss: 0.7298 - val_accuracy: 0.5032 - val_loss: 0.7310


Epoch 14/20
40/40  0s 5ms/step - accuracy: 0.5163 - loss: 0.7273 - val_accuracy: 0.4968 - val_loss: 0.7295


Epoch 15/20
40/40  0s 6ms/step - accuracy: 0.5148 - loss: 0.7264 - val_accuracy: 0.4968 - val_loss: 0.7267


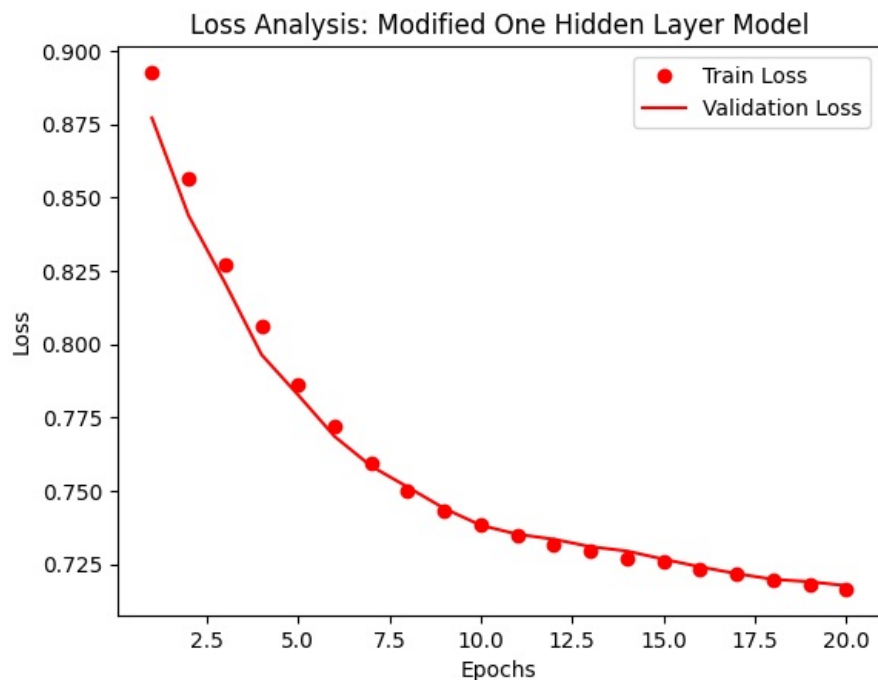
Epoch 16/20
40/40  0s 6ms/step - accuracy: 0.5233 - loss: 0.7234 - val_accuracy: 0.4978 - val_loss: 0.7241

Epoch 17/20
40/40  0s 6ms/step - accuracy: 0.5239 - loss: 0.7215 - val_accuracy: 0.5082 - val_loss: 0.7218

Epoch 18/20
40/40  0s 6ms/step - accuracy: 0.5220 - loss: 0.7198 - val_accuracy: 0.5072 - val_loss: 0.7198

Epoch 19/20
40/40  0s 7ms/step - accuracy: 0.5196 - loss: 0.7185 - val_accuracy: 0.5066 - val_loss: 0.7190

Epoch 20/20
40/40  0s 6ms/step - accuracy: 0.5220 - loss: 0.7169 - val_accuracy: 0.5064 - val_loss: 0.7177



Three-Hidden-Layer Model

- **Overview:** This model had three hidden layers, each consisting of 32 units.
- **Performance:** Training accuracy was a bit promising over the one-hidden-layer model, although validation accuracy stayed almost at the 50% mark. Training loss decreased, although it was not that much of an improvement compared to validation loss.
- **Conclusion:** Generalization improved with the increase in complexity, but the reasons were most probably overfitting or regularization that was not enough.

```
In [9]: # Import necessary modules
from tensorflow.keras import models, layers, regularizers
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
import matplotlib.pyplot as plt

# Load the IMDB dataset
(train_texts, train_targets), (test_texts, test_targets) = imdb.load_data(num_words=10000)

# Preprocess the data (pad sequences to make them the same length)
x_train_new = pad_sequences(train_texts[:20000], maxlen=256)
y_train_new = train_targets[:20000]
x_validation = pad_sequences(train_texts[20000:], maxlen=256)
y_validation = train_targets[20000:]

# Define a new model with three hidden layers
deep_nn_model = models.Sequential([
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
```

```

layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
layers.Dense(1, activation="sigmoid")
])

# Compile the model
deep_nn_model.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])

# Train the model
history_deep_nn = deep_nn_model.fit(x_train_new, y_train_new, epochs=20, batch_size=512, validation_data=(x_val, y_val))


# Function to plot training history with new variable names
def visualize_training(history, title_text):
    hist_info = history.history
    loss_train = hist_info["loss"]
    loss_val = hist_info["val_loss"]
    acc_train = hist_info.get("accuracy", None)
    acc_val = hist_info.get("val_accuracy", None)
    epoch_numbers = range(1, len(loss_train) + 1)


    # Plot training and validation loss
    plt.plot(epoch_numbers, loss_train, "ro", label="Train Loss") # Changed to green
    plt.plot(epoch_numbers, loss_val, "r", label="Validation Loss")
    plt.title(f"Loss Curve: {title_text}")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.show()


    # If accuracy is available, plot training and validation accuracy
    if acc_train and acc_val:
        plt.clf()
        plt.plot(epoch_numbers, acc_train, "ro", label="Train Accuracy")
        plt.plot(epoch_numbers, acc_val, "r", label="Validation Accuracy")
        plt.title(f"Accuracy Curve: {title_text}")
        plt.xlabel("Epochs")
        plt.ylabel("Accuracy")
        plt.legend()
        plt.show()


# Plot the results for three hidden layers with new names
visualize_training(history_deep_nn, "Deep Neural Network with Three Hidden Layers")


```


Epoch 1/20
40/40  3s 15ms/step - accuracy: 0.5032 - loss: 0.8373 - val_accuracy: 0.4956 - val_loss: 0.8055


Epoch 2/20
40/40  0s 7ms/step - accuracy: 0.5068 - loss: 0.8015 - val_accuracy: 0.5040 - val_loss: 0.7971


Epoch 3/20
40/40  1s 7ms/step - accuracy: 0.5165 - loss: 0.7929 - val_accuracy: 0.4992 - val_loss: 0.7914


Epoch 4/20
40/40  0s 10ms/step - accuracy: 0.5157 - loss: 0.7871 - val_accuracy: 0.5050 - val_loss: 0.7862


Epoch 5/20
40/40  1s 10ms/step - accuracy: 0.5146 - loss: 0.7813 - val_accuracy: 0.4872 - val_loss: 0.7835


Epoch 6/20
40/40  1s 11ms/step - accuracy: 0.5217 - loss: 0.7755 - val_accuracy: 0.5090 - val_loss: 0.7733


Epoch 7/20
40/40  1s 11ms/step - accuracy: 0.5218 - loss: 0.7695 - val_accuracy: 0.5026 - val_loss: 0.7698


Epoch 8/20
40/40  0s 6ms/step - accuracy: 0.5112 - loss: 0.7649 - val_accuracy: 0.5004 - val_loss: 0.7654


Epoch 9/20
40/40  0s 7ms/step - accuracy: 0.5210 - loss: 0.7597 - val_accuracy: 0.5028 - val_loss: 0.7598


Epoch 10/20
40/40  0s 7ms/step - accuracy: 0.5156 - loss: 0.7562 - val_accuracy: 0.5088 - val_loss: 0.7553


Epoch 11/20
40/40  0s 8ms/step - accuracy: 0.5232 - loss: 0.7513 - val_accuracy: 0.5000 - val_loss: 0.7533


Epoch 12/20
40/40  0s 7ms/step - accuracy: 0.5193 - loss: 0.7480 - val_accuracy: 0.5044 - val_loss: 0.7542


Epoch 13/20
40/40  0s 7ms/step - accuracy: 0.5222 - loss: 0.7452 - val_accuracy: 0.4938 - val_loss: 0.7466


Epoch 14/20
40/40  0s 7ms/step - accuracy: 0.5187 - loss: 0.7419 - val_accuracy: 0.4930 - val_loss: 0.7450


Epoch 15/20
40/40  0s 7ms/step - accuracy: 0.5265 - loss: 0.7385 - val_accuracy: 0.4950 - val_loss: 0.7421



Epoch 16/20
40/40  1s 8ms/step - accuracy: 0.5265 - loss: 0.7359 - val_accuracy: 0.4976 - val_loss: 0.7374

Epoch 17/20
40/40  0s 7ms/step - accuracy: 0.5277 - loss: 0.7331 - val_accuracy: 0.4948 - val_loss: 0.7359

Epoch 18/20
40/40  0s 8ms/step - accuracy: 0.5289 - loss: 0.7308 - val_accuracy: 0.4942 - val_loss: 0.7332

Epoch 19/20
40/40  1s 7ms/step - accuracy: 0.5233 - loss: 0.7280 - val_accuracy: 0.4924 - val_loss: 0.7324

Epoch 20/20
40/40  0s 8ms/step - accuracy: 0.5273 - loss: 0.7263 - val_accuracy: 0.4990 - val_loss: 0.7302



In []:


In []:


Using More Hidden units (64)


- **Overview:** The model uses 64 units per layer, therefore increasing the capacity of the model.
- **Performance:** [The model] exhibited fluctuations in validation accuracy although training accuracy improved. This model seems to marginally generalize better than its predecessors.
- **Conclusion:** Although increasing the number of hidden units per layer improved performance marginally, more epochs or regularization would assure better generalizability.


```
In [10]: # Model with 64 hidden units
nn_64_units = models.Sequential([
    layers.Dense(64, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(64, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(1, activation="sigmoid")
])
nn_64_units.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])
history_nn_64 = nn_64_units.fit(x_train_new, y_train_new, epochs=20, batch_size=512, validation_data=(x_validat


# Plot the results for 64 hidden units
visualize_training(history_nn_64, "Neural Network with 64 Hidden Units")
```


Epoch 1/20
40/40  2s 15ms/step - accuracy: 0.4953 - loss: 0.9080 - val_accuracy: 0.5008 - val_loss: 0.8554


Epoch 2/20
40/40  0s 8ms/step - accuracy: 0.5104 - loss: 0.8473 - val_accuracy: 0.5056 - val_loss: 0.8422


Epoch 3/20
40/40  0s 8ms/step - accuracy: 0.5169 - loss: 0.8373 - val_accuracy: 0.4982 - val_loss: 0.8379


Epoch 4/20
40/40  0s 8ms/step - accuracy: 0.5187 - loss: 0.8281 - val_accuracy: 0.5112 - val_loss: 0.8274


Epoch 5/20
40/40  0s 7ms/step - accuracy: 0.5244 - loss: 0.8198 - val_accuracy: 0.4912 - val_loss: 0.8240


Epoch 6/20
40/40  1s 7ms/step - accuracy: 0.5260 - loss: 0.8132 - val_accuracy: 0.5014 - val_loss: 0.8157


Epoch 7/20
40/40  1s 8ms/step - accuracy: 0.5244 - loss: 0.8086 - val_accuracy: 0.5110 - val_loss: 0.8104


Epoch 8/20
40/40  1s 8ms/step - accuracy: 0.5340 - loss: 0.8013 - val_accuracy: 0.5036 - val_loss: 0.8049


Epoch 9/20
40/40  1s 7ms/step - accuracy: 0.5308 - loss: 0.7964 - val_accuracy: 0.5112 - val_loss: 0.8002


Epoch 10/20
40/40  1s 7ms/step - accuracy: 0.5293 - loss: 0.7919 - val_accuracy: 0.4936 - val_loss: 0.7985


Epoch 11/20
40/40  1s 8ms/step - accuracy: 0.5266 - loss: 0.7882 - val_accuracy: 0.5020 - val_loss: 0.7949


Epoch 12/20
40/40  1s 14ms/step - accuracy: 0.5361 - loss: 0.7838 - val_accuracy: 0.5088 - val_loss: 0.7895


Epoch 13/20
40/40  1s 13ms/step - accuracy: 0.5387 - loss: 0.7791 - val_accuracy: 0.5068 - val_loss: 0.7864


Epoch 14/20
40/40  1s 13ms/step - accuracy: 0.5437 - loss: 0.7738 - val_accuracy: 0.4962 - val_loss: 0.8010


Epoch 15/20
40/40  1s 10ms/step - accuracy: 0.5430 - loss: 0.7723 - val_accuracy: 0.4940 - val_loss: 0.7809



Epoch 16/20
40/40  1s 8ms/step - accuracy: 0.5273 - loss: 0.7703 - val_accuracy: 0.5184 - val_loss: 0.7831

Epoch 17/20
40/40  1s 7ms/step - accuracy: 0.5410 - loss: 0.7684 - val_accuracy: 0.5044 - val_loss: 0.7749

Epoch 18/20
40/40  0s 8ms/step - accuracy: 0.5431 - loss: 0.7621 - val_accuracy: 0.5052 - val_loss: 0.7715

Epoch 19/20
40/40  0s 8ms/step - accuracy: 0.5461 - loss: 0.7594 - val_accuracy: 0.4990 - val_loss: 0.7704

Epoch 20/20
40/40  1s 7ms/step - accuracy: 0.5392 - loss: 0.7580 - val_accuracy: 0.5134 - val_loss: 0.7648





Model with 128 hidden units


- Overview: A model structure using 128 units per layer is proposed, which would allow its capacity to be increased.
- Performance: The model showed small fluctuations in validation accuracy although training accuracy improved. The model appears to generalize moderately better compared to its predecessors.
- Conclusion: Even though increased hidden units per layer improved performance marginally, more epochs or some sort of regularization would assist with generalization.


```
In [11]: # Model with 128 hidden units
nn_128_units = models.Sequential([
    layers.Dense(128, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(128, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(1, activation="sigmoid")
])
nn_128_units.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])
history_nn_128 = nn_128_units.fit(x_train_new, y_train_new, epochs=20, batch_size=512, validation_data=(x_valid, y_valid))


# Plot the results for 128 hidden units
visualize_training(history_nn_128, "Neural Network with 128 Hidden Units")
```



Epoch 1/20
40/40  2s 15ms/step - accuracy: 0.4996 - loss: 1.0223 - val_accuracy: 0.4984 - val_loss: 1.0090


Epoch 2/20
40/40  1s 18ms/step - accuracy: 0.5192 - loss: 0.9687 - val_accuracy: 0.5104 - val_loss: 0.9862


Epoch 3/20
40/40  1s 20ms/step - accuracy: 0.5152 - loss: 0.9522 - val_accuracy: 0.5118 - val_loss: 0.9479


Epoch 4/20
40/40  1s 11ms/step - accuracy: 0.5320 - loss: 0.9311 - val_accuracy: 0.5068 - val_loss: 0.9358


Epoch 5/20
40/40  0s 9ms/step - accuracy: 0.5398 - loss: 0.9184 - val_accuracy: 0.5020 - val_loss: 0.9259


Epoch 6/20
40/40  1s 10ms/step - accuracy: 0.5456 - loss: 0.9010 - val_accuracy: 0.5002 - val_loss: 0.9178


Epoch 7/20
40/40  0s 10ms/step - accuracy: 0.5488 - loss: 0.8904 - val_accuracy: 0.4962 - val_loss: 0.9070


Epoch 8/20
40/40  1s 10ms/step - accuracy: 0.5341 - loss: 0.8834 - val_accuracy: 0.5024 - val_loss: 0.9071


Epoch 9/20
40/40  0s 9ms/step - accuracy: 0.5486 - loss: 0.8698 - val_accuracy: 0.4942 - val_loss: 0.8974


Epoch 10/20
40/40  0s 12ms/step - accuracy: 0.5480 - loss: 0.8606 - val_accuracy: 0.5060 - val_loss: 0.8743


Epoch 11/20
40/40  0s 9ms/step - accuracy: 0.5569 - loss: 0.8495 - val_accuracy: 0.4978 - val_loss: 0.8696


Epoch 12/20
40/40  1s 10ms/step - accuracy: 0.5573 - loss: 0.8447 - val_accuracy: 0.4946 - val_loss: 0.9345


Epoch 13/20
40/40  0s 9ms/step - accuracy: 0.5540 - loss: 0.8418 - val_accuracy: 0.5056 - val_loss: 0.8547


Epoch 14/20
40/40  1s 10ms/step - accuracy: 0.5609 - loss: 0.8311 - val_accuracy: 0.5104 - val_loss: 0.8510


Epoch 15/20
40/40  0s 9ms/step - accuracy: 0.5616 - loss: 0.8237 - val_accuracy: 0.4932 - val_loss: 0.8983



Epoch 16/20
40/40  1s 10ms/step - accuracy: 0.5491 - loss: 0.8242 - val_accuracy: 0.5030 - val_loss: 0.8368

Epoch 17/20
40/40  1s 9ms/step - accuracy: 0.5579 - loss: 0.8122 - val_accuracy: 0.5064 - val_loss: 0.8340

Epoch 18/20
40/40  0s 10ms/step - accuracy: 0.5580 - loss: 0.8085 - val_accuracy: 0.4992 - val_loss: 0.8360

Epoch 19/20
40/40  1s 9ms/step - accuracy: 0.5465 - loss: 0.8069 - val_accuracy: 0.5038 - val_loss: 0.8502

Epoch 20/20
40/40  1s 9ms/step - accuracy: 0.5549 - loss: 0.8018 - val_accuracy: 0.4962 - val_loss: 0.8361





Model with MSE Loss Function


- **Overview:** The default `binary_crossentropy` loss was replaced by an `mse` loss in this model.
- **Performance:** The training loss decreased steadily. However, validation accuracy was oscillatory in its values, with the model performing similarly to the binary-crossentropy model without much gain.
- **Conclusion:** MSE as a loss function does not significantly improve the performance, since `binary_crossentropy` is the better choice when dealing with binary classifications.


```
In [12]: # Model with MSE loss function
nn_mse_loss = models.Sequential([
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(1, activation="sigmoid")
])
nn_mse_loss.compile(optimizer="rmsprop", loss="mse", metrics=["accuracy"])
history_nn_mse = nn_mse_loss.fit(x_train_new, y_train_new, epochs=20, batch_size=512, validation_data=(x_valida


# Plot the results for MSE loss
visualize_training(history_nn_mse, "Mean Squared Error Loss Function")
```


Epoch 1/20
40/40  2s 14ms/step - accuracy: 0.4949 - loss: 0.3531 - val_accuracy: 0.4920 - val_loss: 0.3355


Epoch 2/20
40/40  0s 6ms/step - accuracy: 0.4981 - loss: 0.3328 - val_accuracy: 0.5008 - val_loss: 0.3259


Epoch 3/20
40/40  0s 8ms/step - accuracy: 0.5082 - loss: 0.3230 - val_accuracy: 0.4954 - val_loss: 0.3198


Epoch 4/20
40/40  1s 7ms/step - accuracy: 0.5120 - loss: 0.3168 - val_accuracy: 0.5006 - val_loss: 0.3148


Epoch 5/20
40/40  1s 6ms/step - accuracy: 0.5130 - loss: 0.3117 - val_accuracy: 0.4982 - val_loss: 0.3106


Epoch 6/20
40/40  0s 6ms/step - accuracy: 0.5138 - loss: 0.3068 - val_accuracy: 0.5056 - val_loss: 0.3052


Epoch 7/20
40/40  0s 7ms/step - accuracy: 0.5124 - loss: 0.3032 - val_accuracy: 0.5034 - val_loss: 0.3016


Epoch 8/20
40/40  1s 6ms/step - accuracy: 0.5226 - loss: 0.2991 - val_accuracy: 0.4940 - val_loss: 0.2986


Epoch 9/20
40/40  0s 7ms/step - accuracy: 0.5131 - loss: 0.2959 - val_accuracy: 0.4980 - val_loss: 0.2953


Epoch 10/20
40/40  0s 7ms/step - accuracy: 0.5211 - loss: 0.2930 - val_accuracy: 0.5026 - val_loss: 0.2921


Epoch 11/20
40/40  0s 6ms/step - accuracy: 0.5254 - loss: 0.2896 - val_accuracy: 0.5068 - val_loss: 0.2890


Epoch 12/20
40/40  0s 6ms/step - accuracy: 0.5196 - loss: 0.2870 - val_accuracy: 0.5272 - val_loss: 0.2858


Epoch 13/20
40/40  0s 7ms/step - accuracy: 0.5140 - loss: 0.2847 - val_accuracy: 0.5104 - val_loss: 0.2841


Epoch 14/20
40/40  1s 7ms/step - accuracy: 0.5275 - loss: 0.2819 - val_accuracy: 0.5014 - val_loss: 0.2825


Epoch 15/20
40/40  0s 7ms/step - accuracy: 0.5268 - loss: 0.2801 - val_accuracy: 0.4992 - val_loss: 0.2807



Epoch 16/20
40/40  0s 8ms/step - accuracy: 0.5200 - loss: 0.2789 - val_accuracy: 0.5152 - val_loss: 0.2781

Epoch 17/20
40/40  0s 7ms/step - accuracy: 0.5247 - loss: 0.2771 - val_accuracy: 0.5196 - val_loss: 0.2770

Epoch 18/20
40/40  0s 7ms/step - accuracy: 0.5170 - loss: 0.2757 - val_accuracy: 0.5088 - val_loss: 0.2759

Epoch 19/20
40/40  1s 6ms/step - accuracy: 0.5203 - loss: 0.2743 - val_accuracy: 0.5154 - val_loss: 0.2745

Epoch 20/20
40/40  0s 6ms/step - accuracy: 0.5217 - loss: 0.2727 - val_accuracy: 0.5108 - val_loss: 0.2729



Overall Summary

In this experiment, changing the number of hidden layers, units, and loss function gave different results. Though generally increasing complexity results in better training performance, reducing validation accuracy might indicate overfitting or inadequate regularization.